

# PIDNet: A Real-time Semantic Segmentation Network Inspired from PID Controller

State of the Art Real-Time Semantic Segmentation on CamVid


State of the Art Real-Time Semantic Segmentation on Cityscapes test

State of the Art Real-Time Semantic Segmentation on Cityscapes val

License MIT

This is the official repository for our recent work: PIDNet ([PDF](#))

## Highlights



Comparison of inference speed and accuracy for real-time models on test set of Cityscapes.


- **Towards Real-time Applications:** PIDNet could be directly used for the real-time applications, such as autonomous vehicle and medical imaging.
- **A Novel Three-branch Network:** Additional boundary branch is introduced to two-branch network to mimic the PID controller architecture and remedy the overshoot issue of previous models.
- **More Accurate and Faster:** PIDNet-S presents 78.6% mIOU with speed of 93.2 FPS on Cityscapes test set and 80.1% mIOU with speed of 153.7 FPS on CamVid test set. Also, PIDNet-L becomes the most accurate one (80.6% mIOU) among all the real-time networks for Cityscapes.

## Updates


- Fixed the data bug for Camvid and the new version of arXiv preprint will be available on Jun 13th. (Jun/09/2022)
- Our paper was marked as state of the art in [Papers with Code](#). (Jun/06/2022)
- Our paper was submitted to arXiv for public access. (Jun/04/2022)
- The training and testing codes and trained models for PIDNet are available here. (Jun/03/2022)

## Demos

A demo of the segmentation performance of our proposed PIDNets: Original video (left) and predictions of PIDNet-S (middle) and PIDNet-L (right)




Cityscapes Stuttgart demo video #1



Cityscapes Stuttgart demo video #2

## Overview



An overview of the basic architecture of our proposed Proportional-Integral-Derivative Network (PIDNet).

P, I and D branches are responsible for detail preservation, context embedding and boundary detection, respectively.

## Detailed Implementation

Stage	Operation			Output
0	Conv $3 \times 3$ , 2, C			256
	Conv $3 \times 3$ , 2, C			
1	$m \times \text{RB}$ , 1, C			256
2	$m \times \text{RB}$ , 2, 2C			128
3	$m \times \text{RB}$ , 1, 2C	$n \times \text{RB}$ , 2, 4C	RB, 1, C\2C	64
	Pag	$\leftarrow \rightarrow$	Add	
4	$m \times \text{RB}$ , 1, 2C	$n \times \text{RB}$ , 2, 8C	RBB\RB, 1, 2C	32
	Pag	$\leftarrow \rightarrow$	Add	
5	RBB, 1, 4C	$2 \times \text{RBB}$ , 2, 16C	RBB, 1, 4C	16
6	$\downarrow$	PPM	$\downarrow$	128
	$\rightarrow$	Bag	$\leftarrow$	
	Conv $3 \times 3$ , 1, 256\128			
	Conv $1 \times 1$ , 1, No. Classes			

Instantiation of the PIDNet for semantic segmentation.

For operation, "OP, N, C" means operation OP with stride of N and the No. output channel is C; Output: output size given input size of 1024;  $m \times \text{RB}$ : m residual basic blocks;  $2 \times \text{RBB}$ : 2 residual bottleneck blocks;  $\text{OP}_1 \backslash \text{OP}_2$ :  $\text{OP}_1$  is used for PIDNet-L while  $\text{OP}_2$  is applied in PIDNet-M and PIDNet-S. (m,n,C) are scheduled to be (2,3,32), (2,3,64) and (3,4,64) for PIDNet-S, PIDNet-M and PIDNet-L, respectively.

## Models

For simple reproduction, we provide the ImageNet pretrained models here.

Model (ImageNet)	PIDNet-S	PIDNet-M	PIDNet-L
Link	<a href="#">download</a>	<a href="#">download</a>	<a href="#">download</a>

Also, the finetuned models on Cityscapes and Camvid are available for direct application in road scene parsing.

Model (Cityscapes)	Val (% mIOU)	Test (% mIOU)	FPS
PIDNet-S	78.8	78.6	93.2
PIDNet-M	79.9	79.8	42.2
PIDNet-L	80.9	80.6	31.1

Model (CamVid)	Val (% mIOU)	Test (% mIOU)	FPS
PIDNet-S	-	80.1	153.7
PIDNet-M	-	82.0	85.6

## Prerequisites

This implementation is based on [HRNet-Semantic-Segmentation](#). Please refer to their repository for installation and dataset preparation. The inference speed is tested on single RTX 3090 using the method introduced by [SwiftNet](#). No third-party acceleration lib is used, so you can try [TensorRT](#) or other approaches for faster speed.

# Usage

## 0. Prepare the dataset

- Download the [Cityscapes](#) and [CamVid](#) datasets and unzip them in `data/cityscapes` and `data/camvid` dirs.
- Check if the paths contained in lists of `data/list` are correct for dataset images.

### **:smiley\_cat: Instruction for preparation of CamVid data (remains discussion) :smiley\_cat:**

- Download the images and annotations from [Kaggle](#), where the resolution of images is 960x720 (original);
- Unzip the data and put all the images and all the colored labels into `data/camvid/images/` and `data/camvid/labels`, respectively;
- Following the split of train, val and test sets used in [SegNet-Tutorial](#), we have generated the dataset lists in `data/list/camvid/`;
- Finished!!! (We have open an issue for everyone who's interested in CamVid to discuss where to download the data and if the split in [SegNet-Tutorial](#) is correct. BTW, do not directly use the split in [Kaggle](#), which is wrong and will lead to unnormail high accuracy. We have revised the CamVid content in the paper and you will see the correct results after its announcement.)

## 1. Training

- Download the ImageNet pretrained models and put them into `pretrained_models/imagenet/` dir.
- For example, train the PIDNet-S on Cityscapes with batch size of 12 on 2 GPUs:

```
python tools/train.py --cfg
configs/cityscapes/pidnet_small_cityscapes.yaml GPUS (0,1)
TRAIN.BATCH_SIZE_PER_GPU 6
```

- Or train the PIDNet-L on Cityscapes using train and val sets simultaneously with batch size of 12 on 4 GPUs:

```
python tools/train.py --cfg
configs/cityscapes/pidnet_large_cityscapes_trainval.yaml GPUS (0,1,2,3)
TRAIN.BATCH_SIZE_PER_GPU 3
```

## 2. Evaluation

- Download the finetuned models for Cityscapes and CamVid and put them into `pretrained_models/cityscapes/` and `pretrained_models/camvid/` dirs, respectively.
- For example, evaluate the PIDNet-S on Cityscapes val set:

```
python tools/eval.py --cfg configs/cityscapes/pidnet_small_cityscapes.yaml
\
```

```
TEST.MODEL_FILE
pretrained_models/cityscapes/PIDNet_S_Cityscapes_val.pt
```

- Or, evaluate the PIDNet-M on CamVid test set:

```
python tools/eval.py --cfg configs/camvid/pidnet_medium_camvid.yaml \
    TEST.MODEL_FILE
pretrained_models/camvid/PIDNet_M_Camvid_Test.pt \
    DATASET.TEST_SET list/camvid/test.lst
```

- Generate the testing results of PIDNet-L on Cityscapes test set:

```
python tools/eval.py --cfg
configs/cityscapes/pidnet_large_cityscapes_trainval.yaml \
    TEST.MODEL_FILE
pretrained_models/cityscapes/PIDNet_L_Cityscapes_test.pt \
    DATASET.TEST_SET list/cityscapes/test.lst
```

### 3. Speed Measurement

- Measure the inference speed of PIDNet-S for Cityscapes:

```
python models/speed/pidnet_speed.py --a 'pidnet-s' --c 19 --r 1024 2048
```

- Measure the inference speed of PIDNet-M for CamVid:

```
python models/speed/pidnet_speed.py --a 'pidnet-m' --c 11 --r 720 960
```

### 4. Custom Inputs

- Put all your images in [samples/](#) and then run the command below using Cityscapes pretrained PIDNet-L for image format of .png:

```
python tools/custom.py --a 'pidnet-l' --p
'../pretrained_models/cityscapes/PIDNet_L_Cityscapes_test.pt' --t '.png'
```

## Citation

If you think this implementation is useful for your work, please cite our paper:

```
@misc{xu2022pidnet,  
      title={PIDNet: A Real-time Semantic Segmentation Network Inspired  
from PID Controller},  
      author={Jiacong Xu and Zixiang Xiong and Shankar P. Bhattacharyya},  
      year={2022},  
      eprint={2206.02066},  
      archivePrefix={arXiv},  
      primaryClass={cs.CV}  
}
```

## Acknowledgement

- Our implementation is modified based on [HRNet-Semantic-Segmentation](#).
- Latency measurement code is borrowed from the [DDRNet](#).
- Thanks for their nice contribution.